

# Application of MI in Industries

Experiment 4

Dhruv Singhal || 500075346 || R177219074 || AIML || B3 || Sem 5

In [2]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

In [3]:

```
#Reading data
data=pd.read_csv("AirQualityUCI.csv")

print(data.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9357 entries, 0 to 9356
Data columns (total 15 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   Date        9357 non-null    object 
 1   Time        9357 non-null    object 
 2   CO(GT)     9357 non-null    float64
 3   PT08.S1(CO) 9357 non-null    int64  
 4   NMHC(GT)   9357 non-null    int64  
 5   C6H6(GT)   9357 non-null    float64
 6   PT08.S2(NMHC) 9357 non-null    int64  
 7   NOx(GT)    9357 non-null    int64  
 8   PT08.S3(NOx) 9357 non-null    int64  
 9   NO2(GT)    9357 non-null    int64  
 10  PT08.S4(NO2) 9357 non-null    int64  
 11  PT08.S5(O3) 9357 non-null    int64  
 12  T          9357 non-null    float64
 13  RH         9357 non-null    float64
```

```
14 AH ..... 9357 non-null float64  
dtypes: float64(5), int64(8), object(2)  
memory usage: 1.1+ MB  
None
```

In [4]: `print(data.describe())`

```
CO(GT) PT08.S1(CO) NMHC(GT) C6H6(GT) PT08.S2(NMHC) \  
count 9357.000000 9357.000000 9357.000000 9357.000000 9357.000000  
mean -34.207524 1048.990061 -159.090093 1.865683 894.595276  
std 77.657170 329.832710 139.789093 41.380206 342.333252  
min -200.000000 -200.000000 -200.000000 -200.000000 -200.000000  
25% 0.600000 921.000000 -200.000000 4.000000 711.000000  
50% 1.500000 1053.000000 -200.000000 7.900000 895.000000  
75% 2.600000 1221.000000 -200.000000 13.600000 1105.000000  
max 11.900000 2040.000000 1189.000000 63.700000 2214.000000  
  
NOx(GT) PT08.S3(NOx) NO2(GT) PT08.S4(NO2) PT08.S5(O3) \  
count 9357.000000 9357.000000 9357.000000 9357.000000 9357.000000  
mean 168.616971 794.990168 58.148873 1391.479641 975.072032  
std 257.433866 321.993552 126.940455 467.210125 456.938184  
min -200.000000 -200.000000 -200.000000 -200.000000 -200.000000  
25% 50.000000 637.000000 53.000000 1185.000000 700.000000  
50% 141.000000 794.000000 96.000000 1446.000000 942.000000  
75% 284.000000 960.000000 133.000000 1662.000000 1255.000000  
max 1479.000000 2683.000000 340.000000 2775.000000 2523.000000  
  
T RH AH  
count 9357.000000 9357.000000 9357.000000  
mean 9.778305 39.485380 -6.837604  
std 43.203623 51.216145 38.976670  
min -200.000000 -200.000000 -200.000000  
25% 10.900000 34.100000 0.692300  
50% 17.200000 48.600000 0.976800  
75% 24.100000 61.900000 1.296200  
max 44.600000 88.700000 2.231000
```

In [5]: `print(data.isnull().sum())`

```
Date ..... 0  
Time ..... 0  
CO(GT) ..... 0  
PT08.S1(CO) ..... 0
```

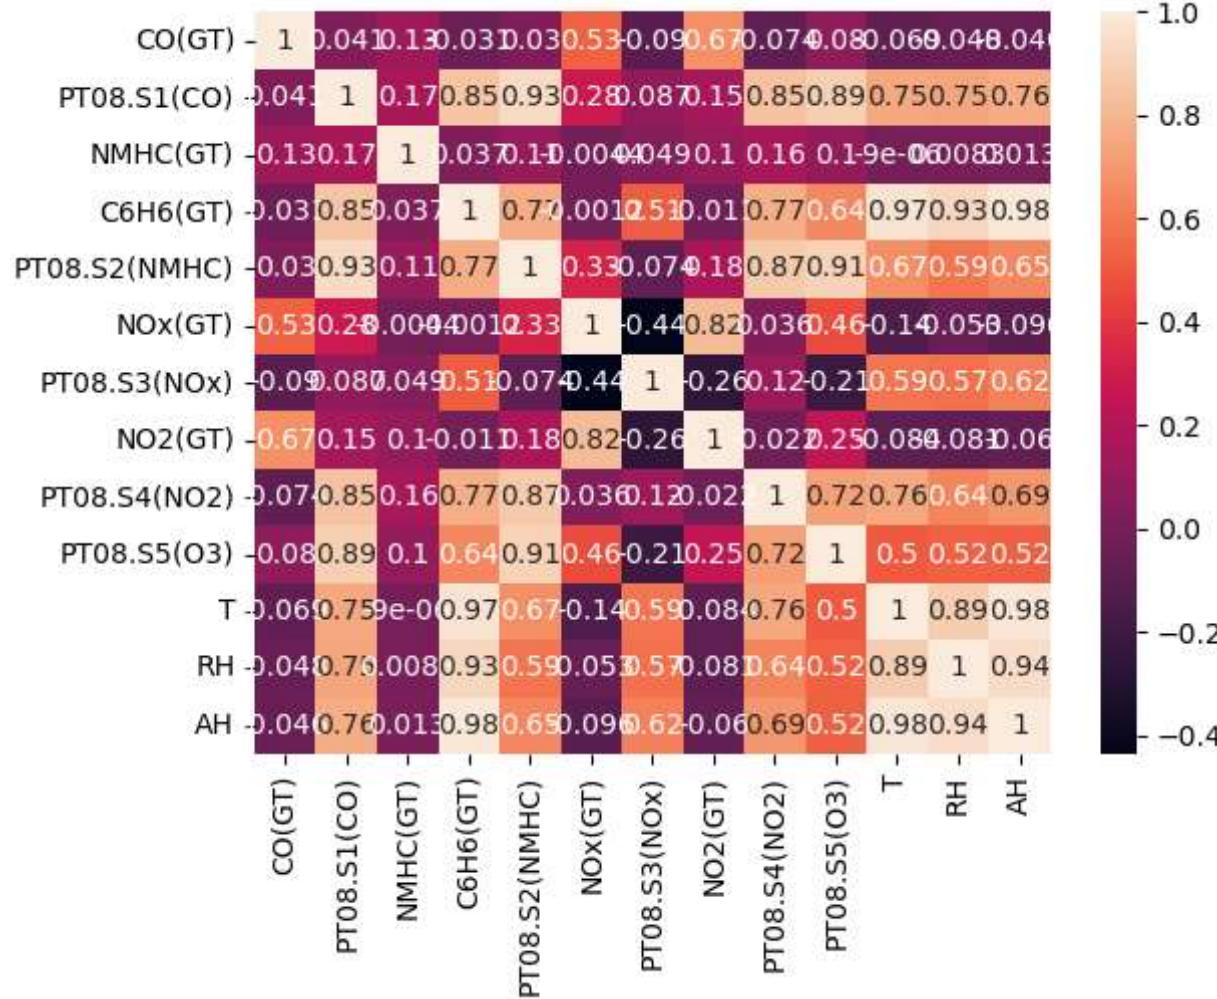
```
NMHC(GT)      0  
C6H6(GT)      0  
PT08.S2(NMHC) 0  
NOx(GT)       0  
PT08.S3(NOx)  0  
NO2(GT)       0  
PT08.S4(NO2)  0  
PT08.S5(O3)   0  
T              0  
RH             0  
AH             0  
dtype: int64
```

```
In [6]: print(data.corr())
```

	CO(GT)	PT08.S1(CO)	NMHC(GT)	C6H6(GT)	PT08.S2(NMHC)	\
CO(GT)	1.000000	0.041411	0.128351	-0.031378	0.029926	
PT08.S1(CO)	0.041411	1.000000	0.170007	0.852687	0.933102	
NMHC(GT)	0.128351	0.170007	1.000000	0.037323	0.110104	
C6H6(GT)	-0.031378	0.852687	0.037323	1.000000	0.767433	
PT08.S2(NMHC)	0.029926	0.933102	0.110104	0.767433	1.000000	
NOx(GT)	0.526451	0.277993	-0.004427	-0.001174	0.331272	
PT08.S3(NOx)	-0.089981	0.087019	0.048821	0.512193	-0.073667	
NO2(GT)	0.671127	0.154030	0.103307	-0.010992	0.176488	
PT08.S4(NO2)	-0.073724	0.845149	0.162680	0.774673	0.874782	
PT08.S5(O3)	0.080310	0.892434	0.101185	0.641334	0.909905	
T	-0.068939	0.754844	-0.000009	0.971375	0.669025	
RH	-0.048227	0.745375	0.008284	0.925062	0.585803	
AH	-0.045892	0.764903	0.012500	0.984555	0.646572	
	NOx(GT)	PT08.S3(NOx)	NO2(GT)	PT08.S4(NO2)	PT08.S5(O3)	\
CO(GT)	0.526451	-0.089981	0.671127	-0.073724	0.080310	
PT08.S1(CO)	0.277993	0.087019	0.154030	0.845149	0.892434	
NMHC(GT)	-0.004427	0.048821	0.103307	0.162680	0.101185	
C6H6(GT)	-0.001174	0.512193	-0.010992	0.774673	0.641334	
PT08.S2(NMHC)	0.331272	-0.073667	0.176488	0.874782	0.909905	
NOx(GT)	1.000000	-0.436084	0.817139	0.035546	0.461889	
PT08.S3(NOx)	-0.436084	1.000000	-0.256232	0.122734	-0.208865	
NO2(GT)	0.817139	-0.256232	1.000000	-0.022174	0.253439	
PT08.S4(NO2)	0.035546	0.122734	-0.022174	1.000000	0.723690	
PT08.S5(O3)	0.461889	-0.208865	0.253439	0.723690	1.000000	
T	-0.138452	0.588111	-0.084104	0.755060	0.503700	
RH	-0.053009	0.573549	-0.081305	0.640707	0.524955	
AH	-0.095847	0.621618	-0.060440	0.691913	0.519467	

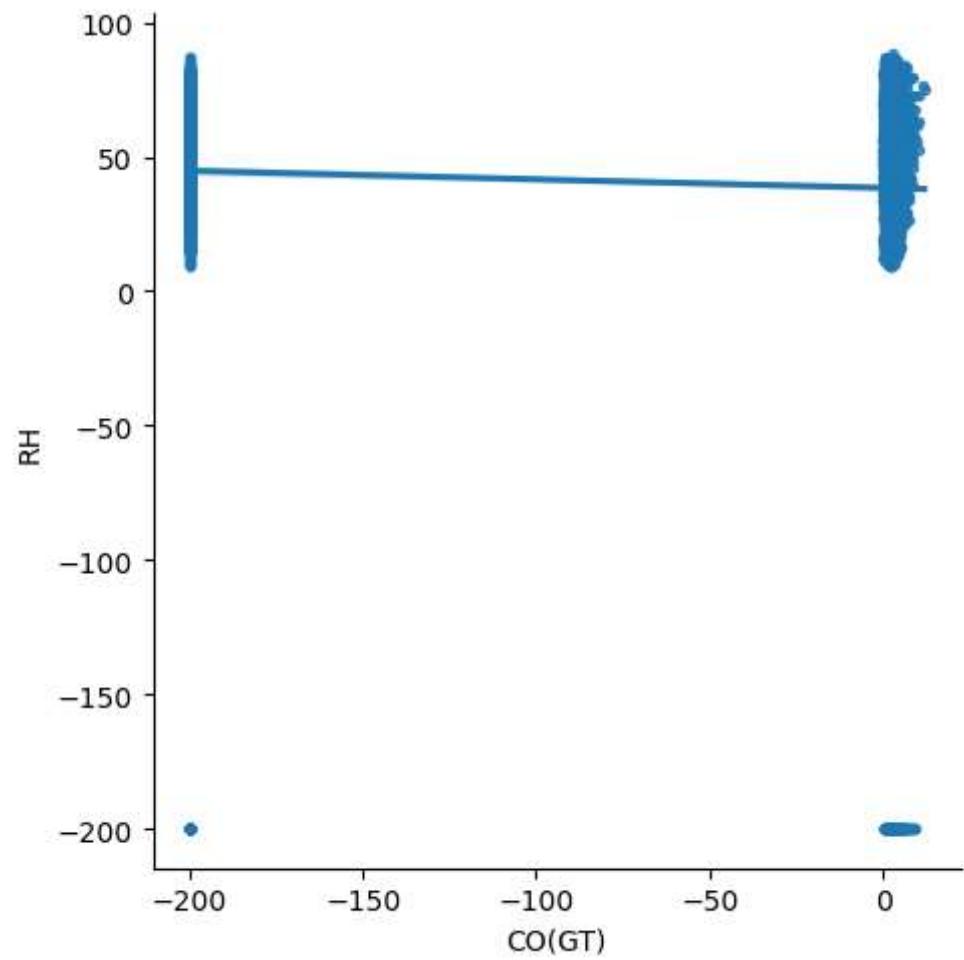
	T	RH	AH
CO(GT)	-0.068939	-0.048227	-0.045892
PT08.S1(CO)	0.754844	0.745375	0.764903
NMHC(GT)	-0.000009	0.008284	0.012500
C6H6(GT)	0.971375	0.925062	0.984555
PT08.S2(NMHC)	0.669025	0.585803	0.646572
NOx(GT)	-0.138452	-0.053009	-0.095847
PT08.S3(NOx)	0.588111	0.573549	0.621618
NO2(GT)	-0.084104	-0.081305	-0.060440
PT08.S4(NO2)	0.755060	0.640707	0.691913
PT08.S5(O3)	0.503700	0.524955	0.519467
T	1.000000	0.885911	0.981001
RH	0.885911	1.000000	0.943995
AH	0.981001	0.943995	1.000000

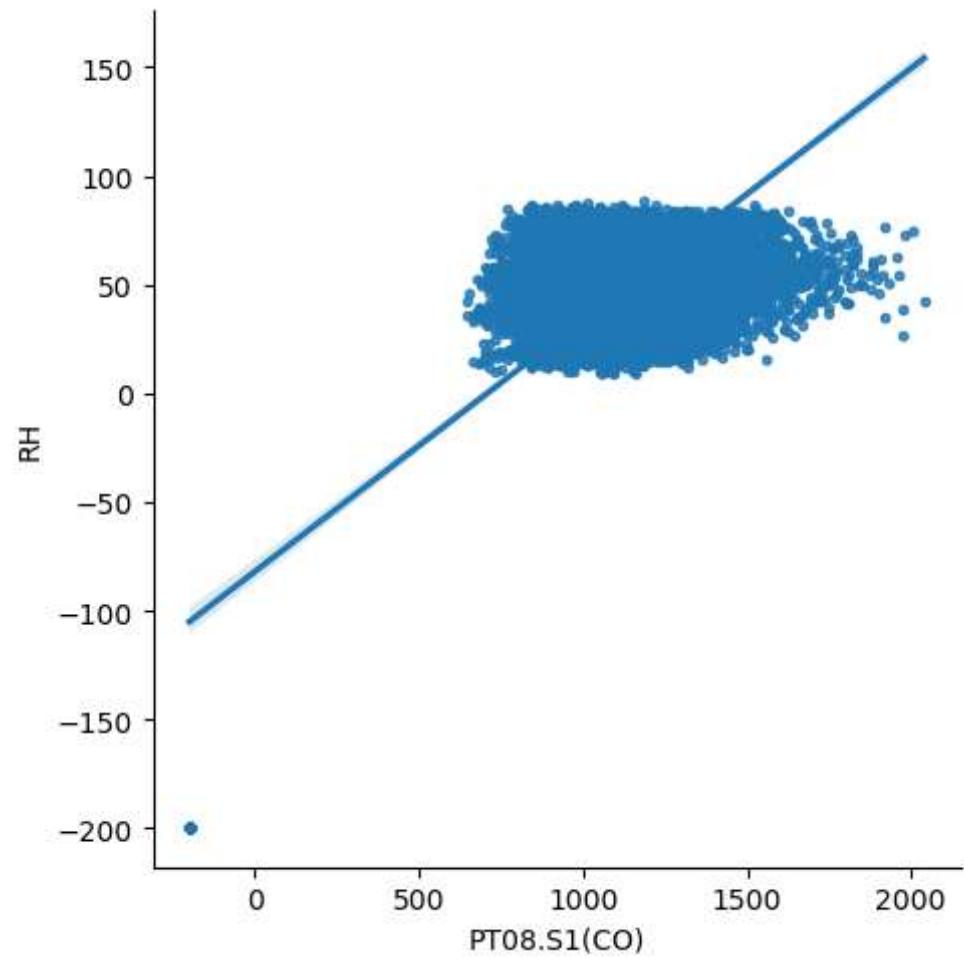
```
In [7]: sns.heatmap(data.corr(), annot=True)
plt.show()
```

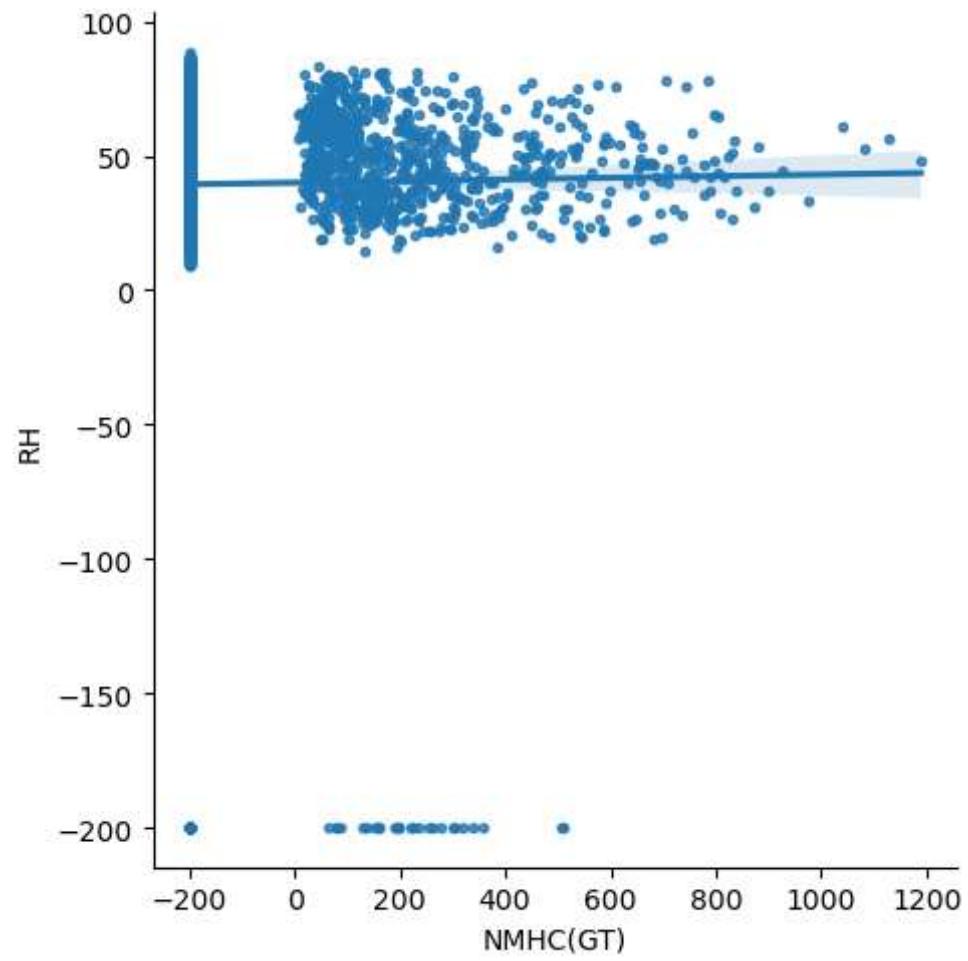


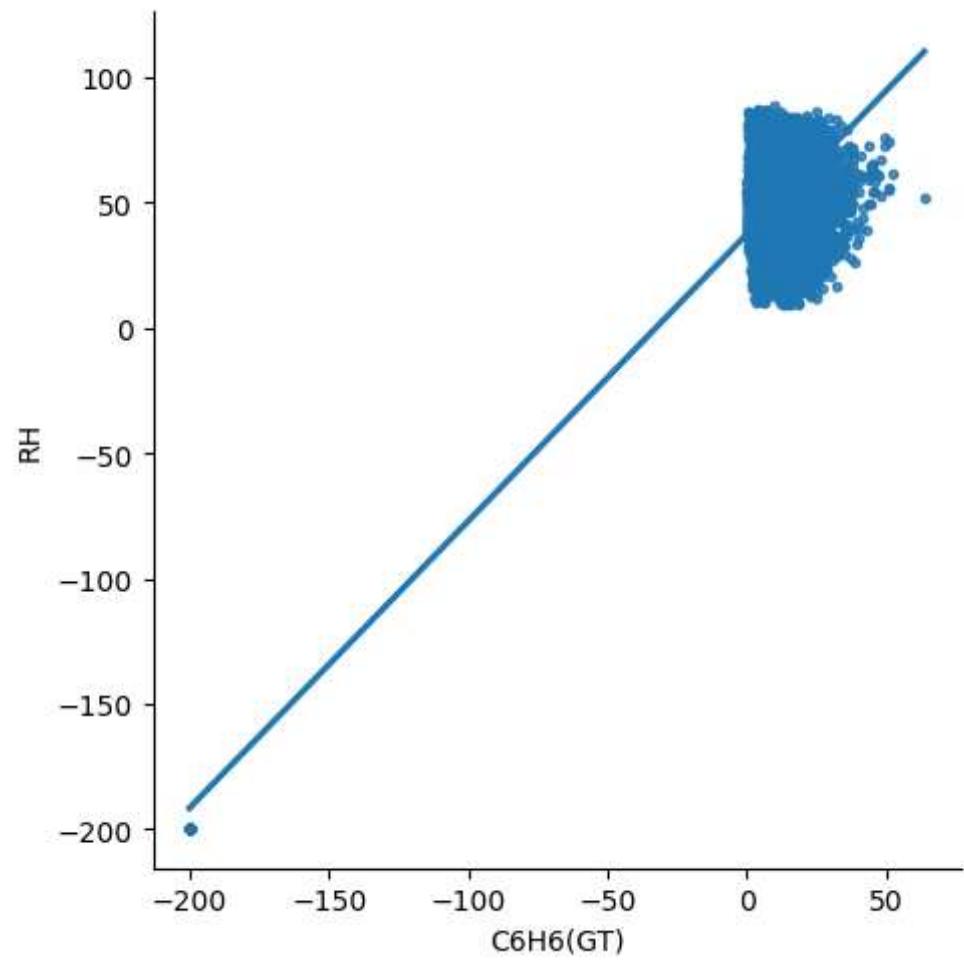
In [8]:

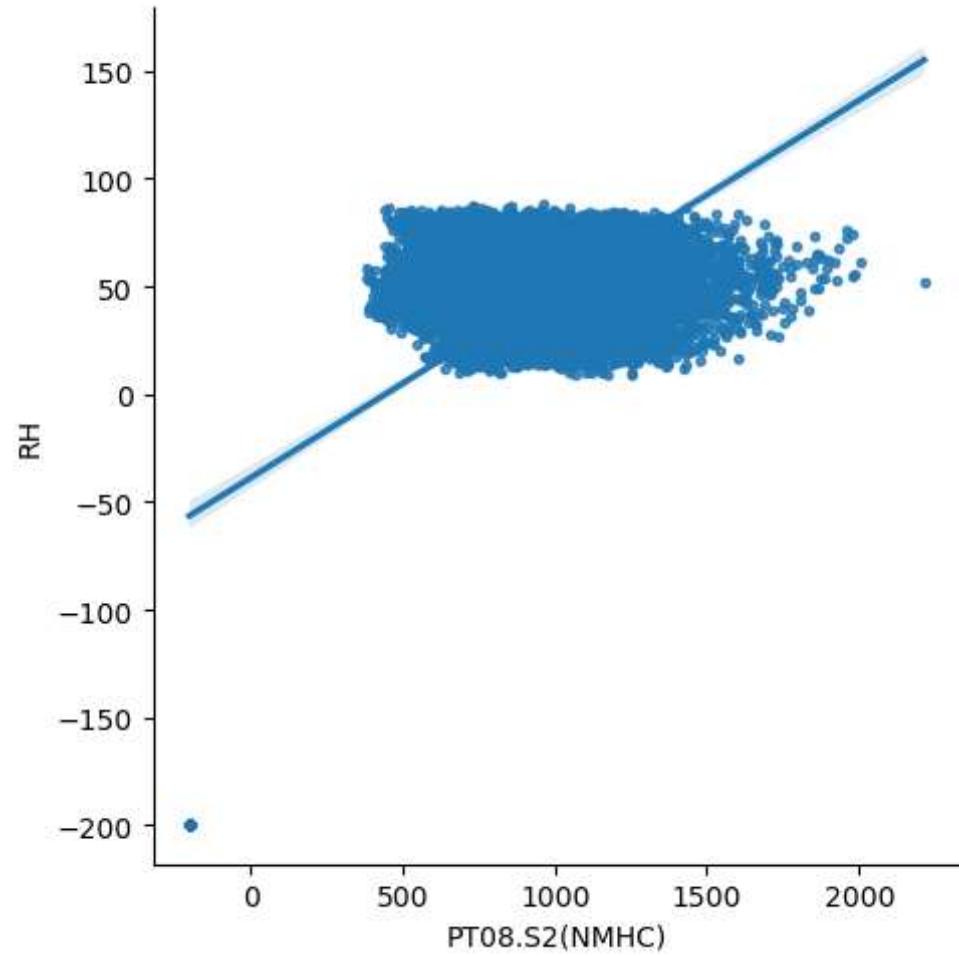
```
col_=data.columns.tolist()[2:]
for i in data.columns.tolist()[2:]:
    sns.lmplot(x=i,y='RH',data=data,markers='.')
plt.show()
```

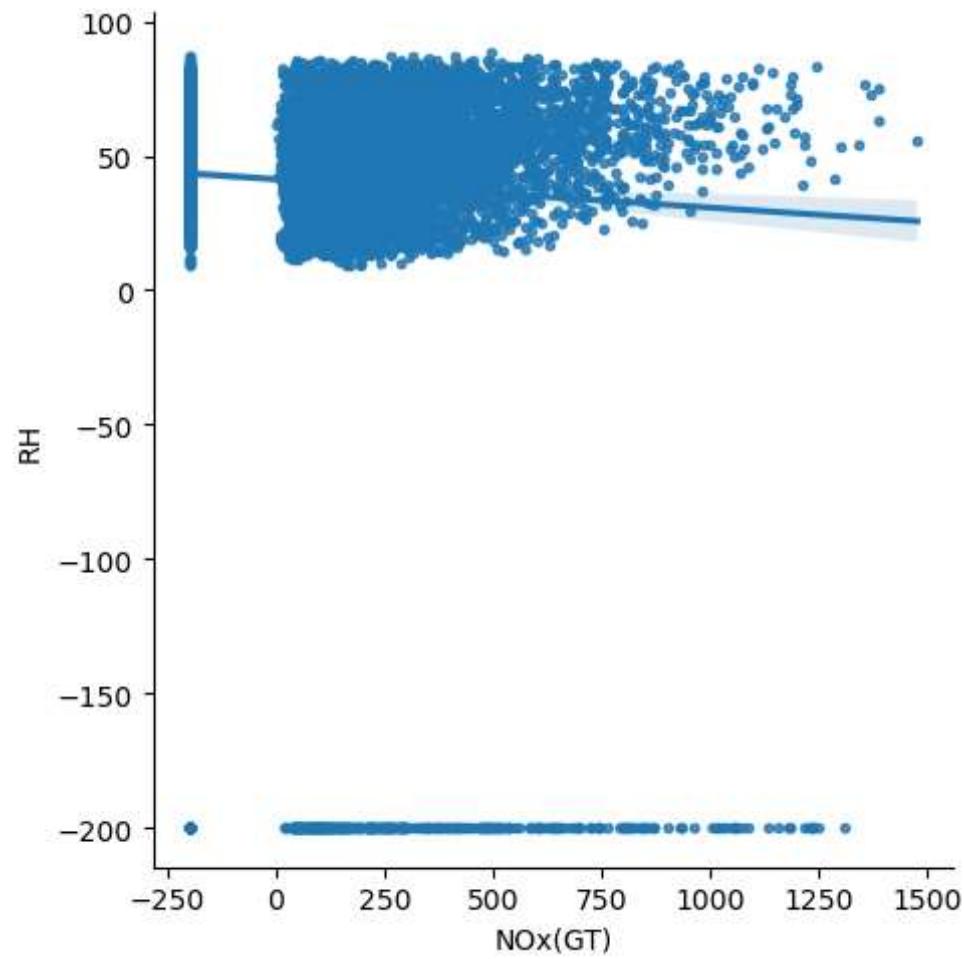


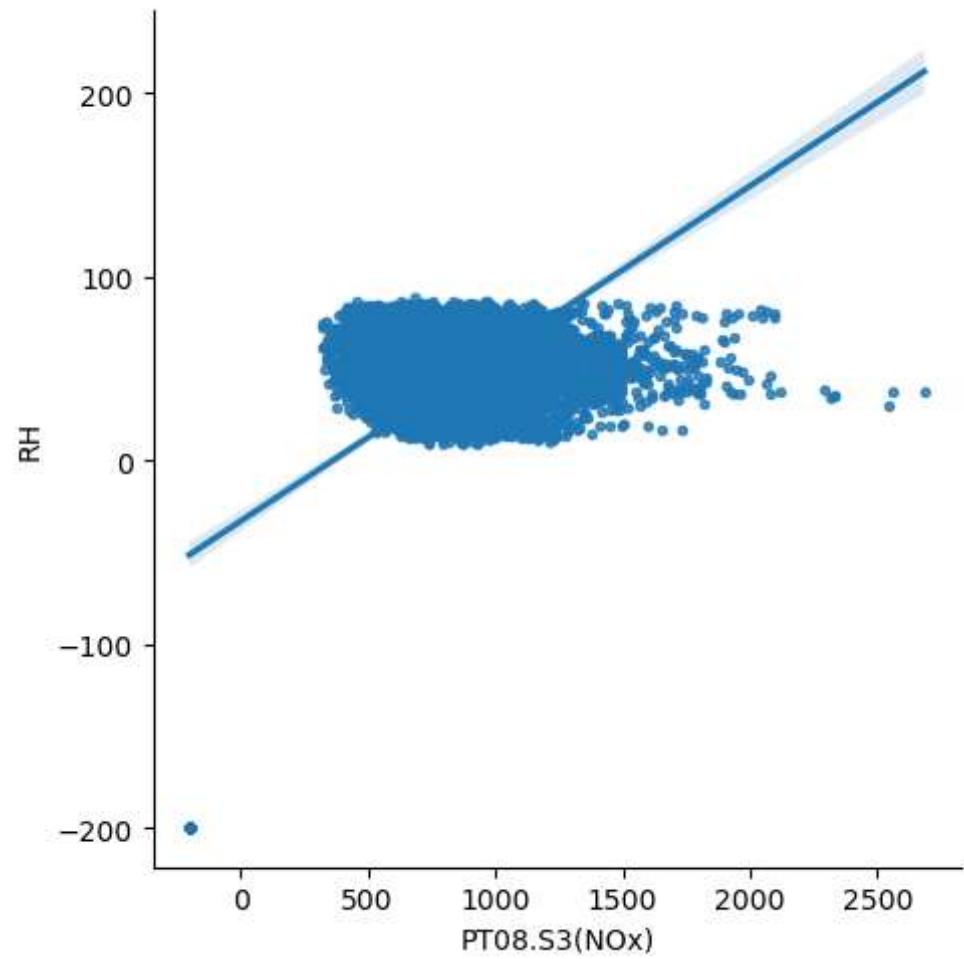


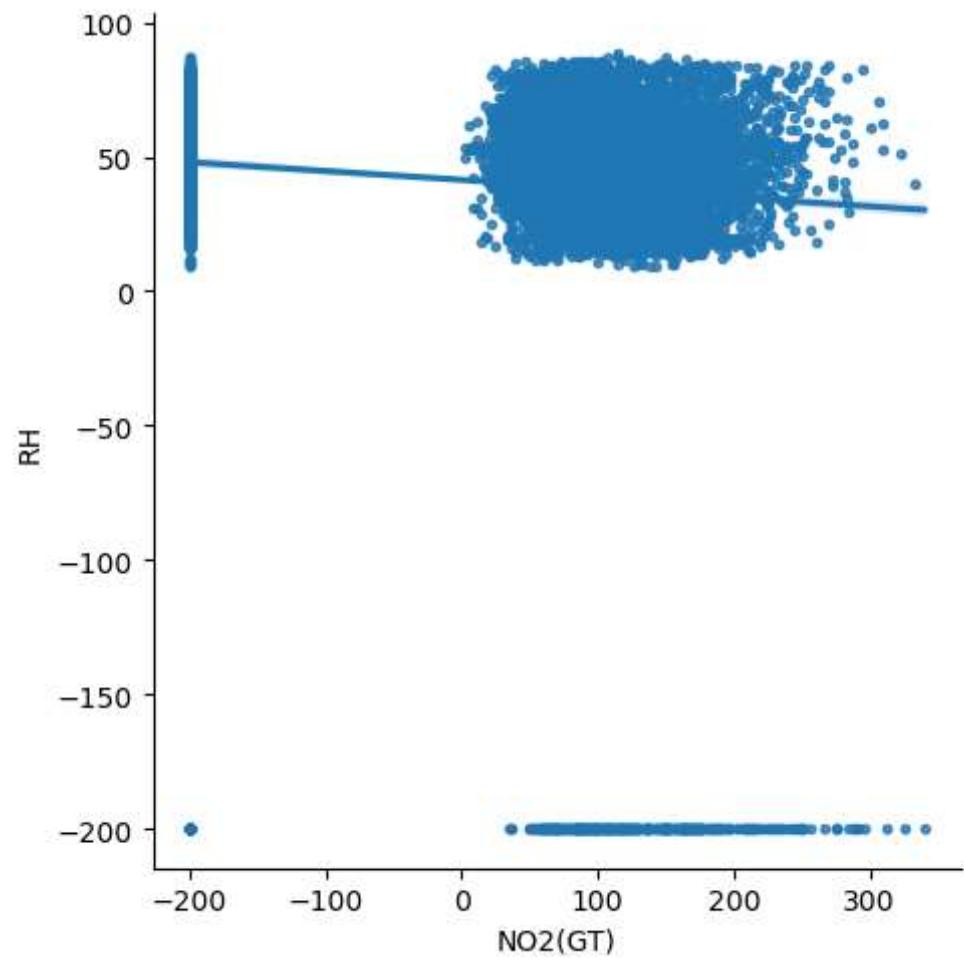


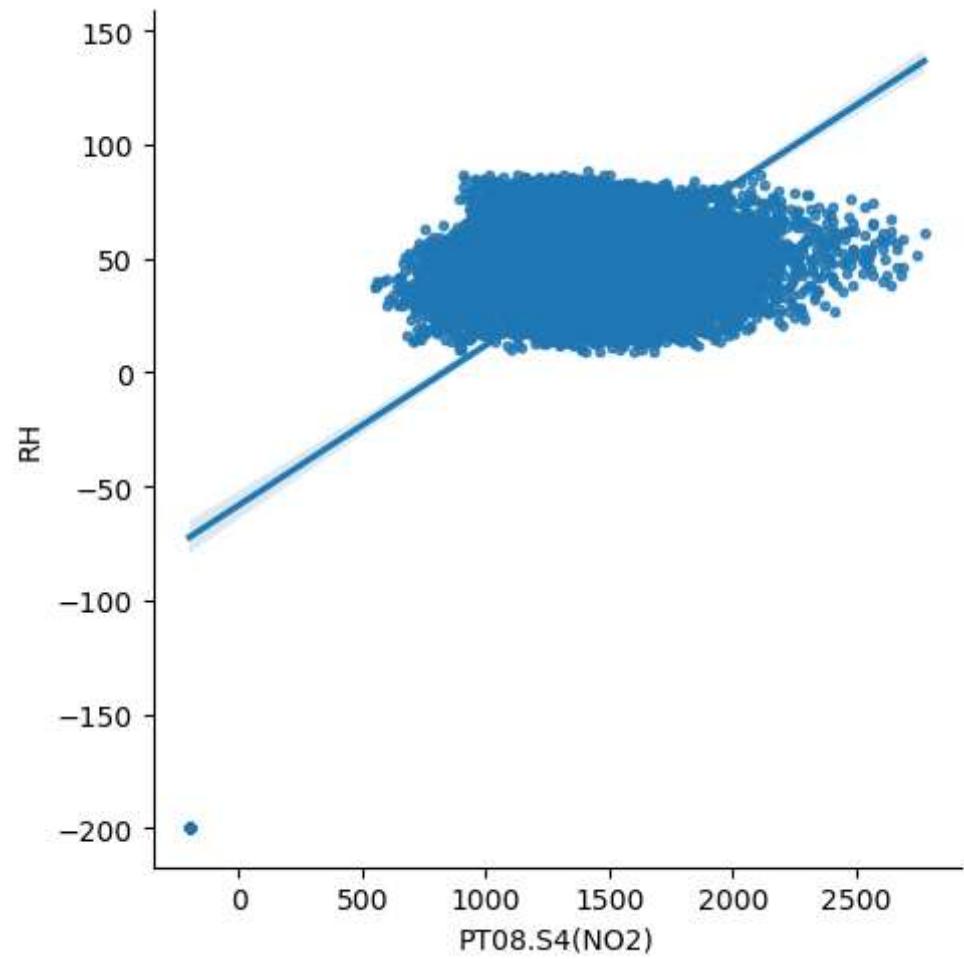


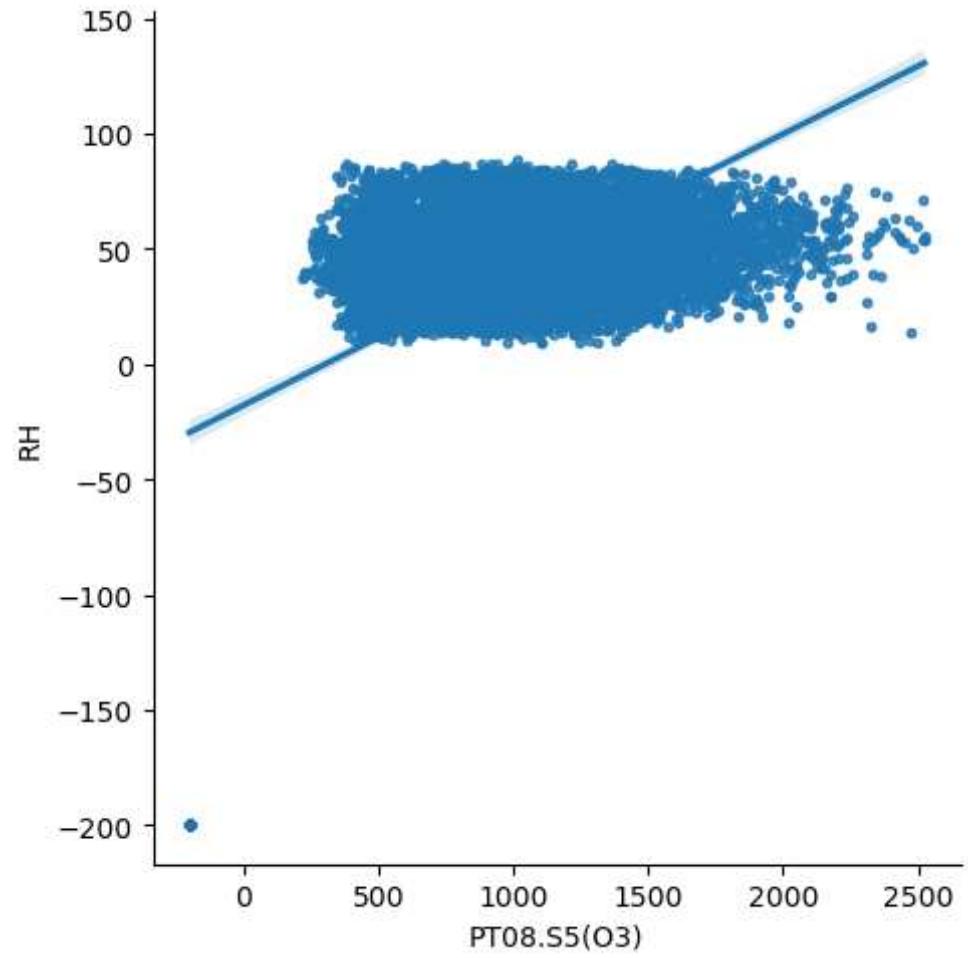


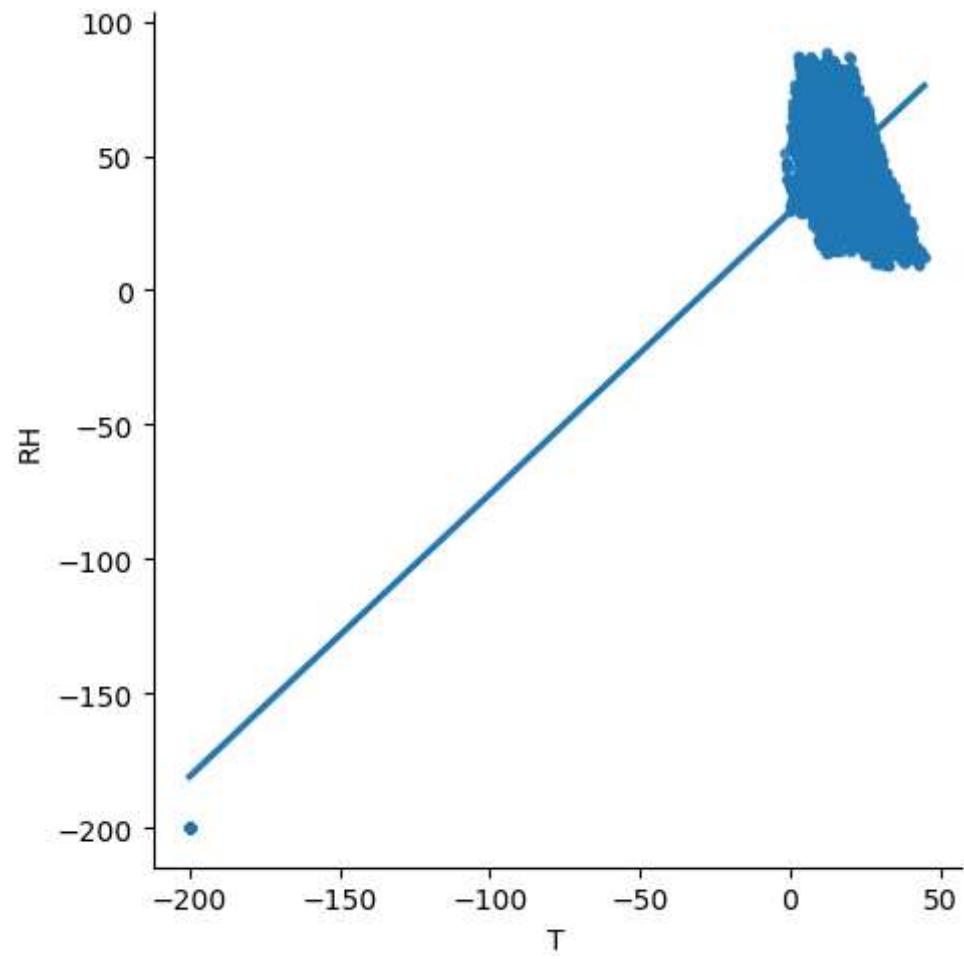


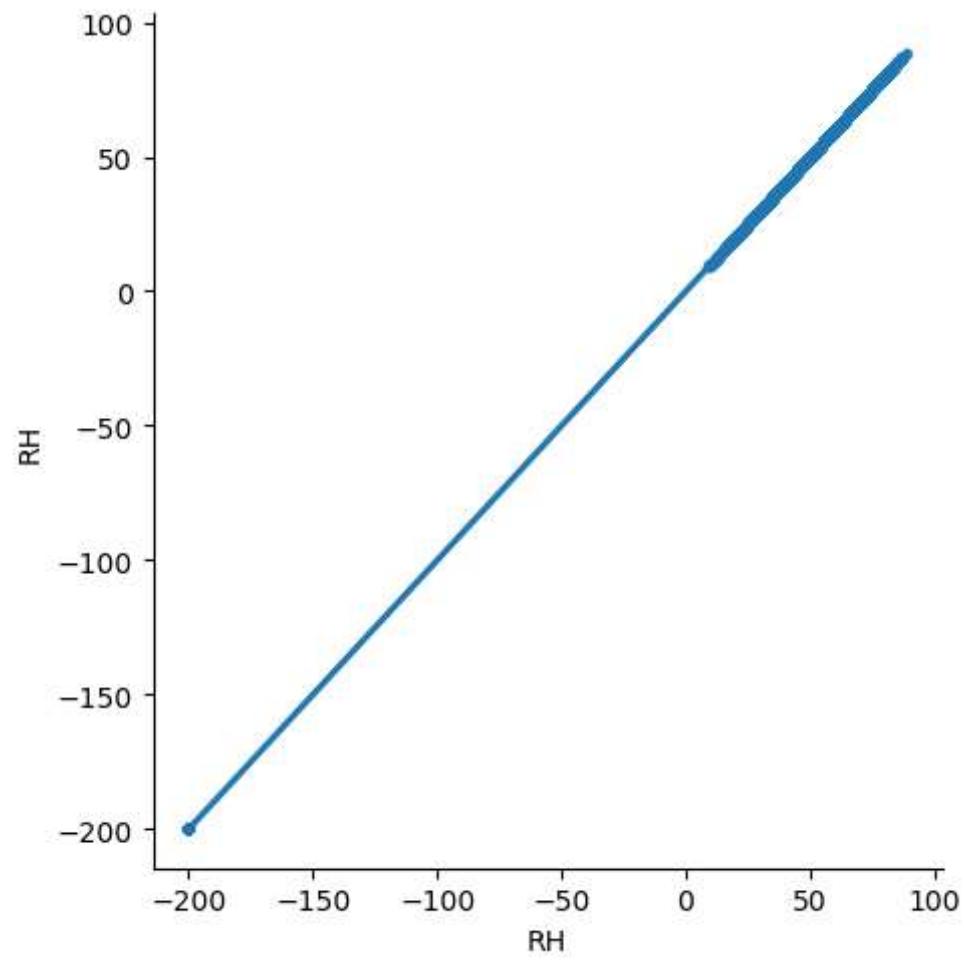


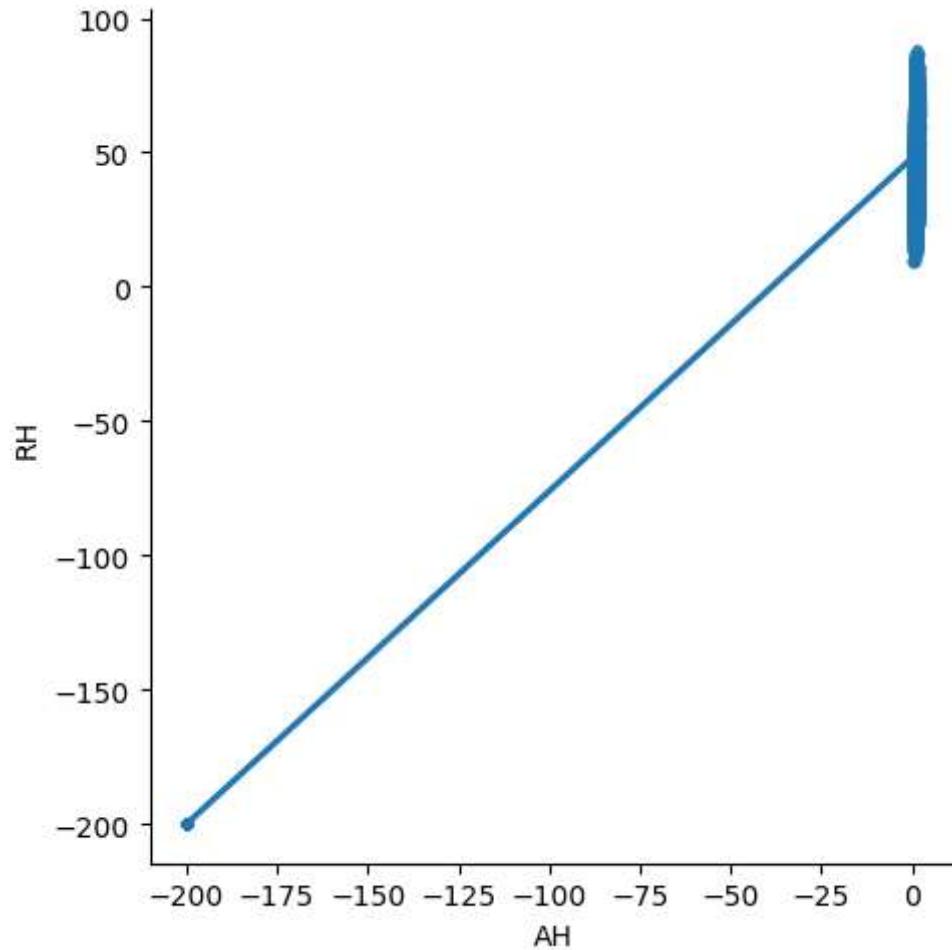












In [9]:

```
from sklearn.preprocessing import StandardScaler      #import normalisation package
from sklearn.model_selection import train_test_split    #import train test split
from sklearn.linear_model import LinearRegression     #import linear regression package
from sklearn.metrics import mean_squared_error,mean_absolute_error  #import mean squared error and mean absolute error
```

In [10]:

```
X=data[col_].drop('RH',1)      #X-input features
y=data['RH']
```

In [11]:

```
ss=StandardScaler()      #initialise
```

```
X_std=ss.fit_transform(X)      #apply standardisation
```

```
In [12]: X_train, X_test, y_train, y_test=train_test_split(X_std,y,test_size=0.3, random_state=42)
```

```
In [13]: print('Training data size:',X_train.shape)
print('Test data size:',X_test.shape)
```

Training data size: (6549, 12)

Test data size: (2808, 12)

```
In [14]: lr=LinearRegression()
lr_model=lr.fit(X_train,y_train)      #fit the linear model on train data
print('Intercept:',lr_model.intercept_)
print('-----')
print('Slope:')
print(list(zip(X.columns.tolist(),lr_model.coef_)))
```

Intercept: 39.48586056471033

-----  
Slope:

```
[('CO(GT)', 0.37631295673257303), ('PT08.S1(CO)', 3.2761957157732318), ('NMHC(GT)', -2.1453360996088615), ('C6H6(GT)', -48.10277307336059), ('PT08.S2(NMHC)', -16.703868112073955), ('NOx(GT)', 7.8025418758624285), ('PT08.S3(Nox)', -6.431116299980261), ('NO2(GT)', -6.404494232464043), ('PT08.S4(NO2)', 27.603146537409394), ('PT08.S5(O3)', -0.16832955987058656), ('T', -87.61049424500278), ('AH', 175.28393938691548)]
```

```
In [15]: y_pred=lr_model.predict(X_test)          #predict using the model
rmse=np.sqrt(mean_squared_error(y_test,y_pred))    #calculate rmse
print('RMSE of model:',rmse)
print("linear regression: ", lr_model.score(X_test, y_test))
```

RMSE of model: 7.879527482532906

linear regression: 0.9762585204595909

```
In [16]: from sklearn.ensemble import RandomForestRegressor      #import random forest regressor
rf_reg=RandomForestRegressor()
```

```
In [17]: rf_model=rf_reg.fit(X_train,y_train)      #fit model
y_pred_rf=rf_model.predict(X_test)            #predict
```

```
In [18]:  
print('RMSE of predicted RH in RF model:',np.sqrt(mean_squared_error(y_test,y_pred_rf)))  
print("Random Forest: ", rf_model.score(X_test, y_test))
```

RMSE of predicted RH in RF model: 0.6025343247779432

Random Forest: 0.9998611739342902

```
In [ ]:
```

```
In [ ]:
```